
Lecture 5

Y86-64 Instruction Set Architecture

Having explored how real programs are represented in x64-64 instructions, in this lecture we will begin to explore how the actual instruction set architecture is implemented in hardware. To do so, we study a subset of the full x86-64 instruction set known as Y86-64. This will allow us to explore the actual representation of instructions and how the hardware control logic actually implements a processor design.

Y86-64 Programmer Visible State

The Y86-64 programmer visible state includes 15 general purpose 64-bit registers, named `%rax`, `%rcx`, `%rdx`, `%rbx`, `%rsp`, `%rbp`, `%rsi`, `%rdi`, `%r8`, through to `%r14`. `%rsp` is the stack pointer used by `push`, `pop`, `call` and `ret` instructions. An additional program counter, `PC`, keeps track of the address of the current instruction in memory.

There are three single-bit condition codes, `ZF`, `SF`, and `OF`, which are updated by arithmetic instructions, and indicate that the most recent operation resulted in a zero value, a negative value, and/or a two's complement overflow.

The memory where both programs and data are stored is represented as an array of bytes, referenced by a virtual address (the hardware and operating system turn this into a real physical address).

Y86-64 Instructions and Encoding

The Y86-64 instruction set is a subset of x86-64. It includes only 64-bit integer operations, and operates on 64-bit data which are referred to as words. The assembly code is similar to that of the x86-64 instruction set, with separate instructions replacing some of the complex operand options of the x86-64 instructions.

- `halt` stops instruction execution.
- `nop` is 'no operation' and does nothing.
- `rrmovq`, `irmovq`, `rmmovq`, `rmovq` are equivalent to the x86-64 `movq` instruction, each taking either a register, immediate value, and memory location as a source, and a register or memory location as a destination.
- `OPq` represents the arithmetical and logical operators `addq`, `subq`, `andq` and `xorq`.
- `JXX` represents a family of jump instructions which use the condition codes to provide `jmp`, `jle`, `jg`, `jle`, `jge`, `jle` and `jg` instructions, all of which take a 64-bit absolute destination.
- `cmovXX` represents a family of conditional move instructions, again using the condition codes to provide `cmovle`, `cmovl`, `cmovbe`, `cmovb`, `cmovge` and `cmovg` instructions. Each indicates a register as the source, and a register as the destination.

Bytes	0	1	2	3	4	5	6	7	8	9
halt	0	0								
nop	1	0								
rrmovq rA,rB	2	0	rA	rB						
irmovq V,rB	3	0	F	rB			V			
rmmovq rA,D(rB)	4	0	rA	rB			D			
rrmovq D(rB),rA	5	0	rA	rB			D			
OPq rA,rB	6	fn	rA	rB						
jXX Dest	7	fn					Dest			
cmovXX rA,rB	2	fn	rA	rB						
call Dest	8	0					Dest			
ret	9	0								
pushq rA	A	0	rA	F						
popq rA	B	0	rA	F						

Table 1: Y86-64 instruction set encoding.

- **call** and **ret** interact with the `%rsp` pointer in the same way as the x86-64 instruction do, with **call** using a 64-bit absolute destination.

Instructions have encoded lengths with varies from a single byte to ten bytes, as shown in Table 1. All addresses and immediate values are 8 bytes long and are encoded as little endian values. Note that many of these instructions are parameterised; using the least significant four bits of the first byte to indicate their operation (see Table 2). Also, many instructions have a *register specifier byte* indicating which registers the instruction operates on. When only one register is required, the unused one is replaced with `0xF`. The coding of the register specifier byte is shown in Table 3.

Y86-64 Exceptions

The state of operation of the processor implementing the Y86-64 instructions is indicated by a status code which can take the values shown in Table 4.

		jmp	<table><tr><td>7</td><td>0</td></tr></table>	7	0	rrmovq	<table><tr><td>2</td><td>0</td></tr></table>	2	0		
7	0										
2	0										
		jle	<table><tr><td>7</td><td>1</td></tr></table>	7	1	cmovle	<table><tr><td>2</td><td>1</td></tr></table>	2	1		
7	1										
2	1										
		jl	<table><tr><td>7</td><td>2</td></tr></table>	7	2	cmovl	<table><tr><td>2</td><td>2</td></tr></table>	2	2		
7	2										
2	2										
addq	<table><tr><td>6</td><td>0</td></tr></table>	6	0	je	<table><tr><td>7</td><td>3</td></tr></table>	7	3	cmove	<table><tr><td>2</td><td>3</td></tr></table>	2	3
6	0										
7	3										
2	3										
subq	<table><tr><td>6</td><td>1</td></tr></table>	6	1	jne	<table><tr><td>7</td><td>4</td></tr></table>	7	4	cmovne	<table><tr><td>2</td><td>4</td></tr></table>	2	4
6	1										
7	4										
2	4										
andq	<table><tr><td>6</td><td>2</td></tr></table>	6	2	jge	<table><tr><td>7</td><td>5</td></tr></table>	7	5	cmovge	<table><tr><td>2</td><td>5</td></tr></table>	2	5
6	2										
7	5										
2	5										
xorq	<table><tr><td>6</td><td>3</td></tr></table>	6	3	jg	<table><tr><td>7</td><td>6</td></tr></table>	7	6	cmovg	<table><tr><td>2</td><td>6</td></tr></table>	2	6
6	3										
7	6										
2	6										

Table 2: Function codes for **OPq**, **jXX** and **cmovXX** instructions.

Table 3: Y86-64 registers.

Number	Register Name	Number	Register Name
0	%rax	8	%r8
1	%rcx	9	%r9
2	%rdx	A	%r10
3	%rbx	B	%r11
4	%rsp	C	%r12
5	%rbp	D	%r13
6	%rsi	E	%r14
7	%rdi	F	No register

Table 4: Y86-64 status codes.

Value	Name	Meaning
1	AOK	Normal operation
2	HLT	halt instruction encountered
3	ADR	Invalid address encountered
4	INS	Invalid instruction encountered

Y86-64 Code Examples

We can simulate the execution of Y86-64 instructions in a number of different ways. In order to fully explore the hardware implementation of the Y86-64 later in the course we will use a simulation of the actual processor hardware. However, here we use an instruction level simulator that can compile Y86-64 instructions to an object file, and then execute the instructions in this file.

Below is an example of Y86-64 code for such a simulator, which implements the addition of the elements of an array (as we saw earlier in C code).

```
.pos 0
    irmovq stack, %rsp
    call main
    halt

.align 8
array:
    .quad 0x0000000000000001
    .quad 0x0000000000000002
    .quad 0x0000000000000003
    .quad 0x0000000000000004

main:
    irmovq array, %rdi
    irmovq $4, %rsi
    call sum
    ret

sum:
    irmovq $8, %r8
    irmovq $1, %r9
    xorq %rax, %rax
    andq %rsi, %rsi
    jmp test
loop:
    mrmovq (%rdi), %r10
    addq %r10, %rax
    addq %r8, %rdi
    subq %r9, %rsi
test:
    jne loop
    ret

.pos 0x200
stack:
```

This code results in the following compiled code when using the **yas** assembler to give as output.

```

0x000:                                | .pos 0
0x000: 30f400020000000000000000 |      irmovq stack, %rsp
0x00a: 803800000000000000000000 |      call main
0x013: 00                             |      halt
                                |
0x018:                                | .align 8
0x018:                                | array:
0x018: 010000000000000000000000 |      .quad 0x000000000000000001
0x020: 020000000000000000000000 |      .quad 0x000000000000000002
0x028: 030000000000000000000000 |      .quad 0x000000000000000003
0x030: 040000000000000000000000 |      .quad 0x000000000000000004
                                |
0x038:                                | main:
0x038: 30f718000000000000000000 |      irmovq array, %rdi
0x042: 30f604000000000000000000 |      irmovq $4, %rsi
0x04c: 805600000000000000000000 |      call sum
0x055: 90                             |      ret
                                |
0x056:                                | sum:
0x056: 30f808000000000000000000 |      irmovq $8, %r8
0x060: 30f901000000000000000000 |      irmovq $1, %r9
0x06a: 6300                          |      xorq %rax, %rax
0x06c: 6266                          |      andq %rsi, %rsi
0x06e: 708700000000000000000000 |      jmp test
0x077:                                | loop:
0x077: 50a700000000000000000000 |      mrmovq (%rdi), %r10
0x081: 60a0                          |      addq %r10, %rax
0x083: 6087                          |      addq %r8, %rdi
0x085: 6196                          |      subq %r9, %rsi
0x087:                                | test:
0x087: 747700000000000000000000 |      jne loop
0x090: 90                             |      ret
                                |
0x200:                                | .pos 0x200
0x200:                                | stack:

```

When run with the `yis` simulator, the following output is generated.

```
Stopped in 34 steps at PC = 0x13.  Status 'HLT', CC Z=1 S=0 O=0
Changes to registers:
%rax:  0x0000000000000000      0x000000000000000a
%rsp:  0x0000000000000000      0x0000000000000200
%rdi:  0x0000000000000000      0x0000000000000038
%r8:   0x0000000000000000      0x0000000000000008
%r9:   0x0000000000000000      0x0000000000000001
%r10:  0x0000000000000000      0x0000000000000004

Changes to memory:
0x01f0: 0x0000000000000000      0x0000000000000055
0x01f8: 0x0000000000000000      0x0000000000000013
```

The output summarises the changes to registers and memory contents. In this case, the changes to the memory correspond to the return addresses pushed to the stack.

Simulator

The `yas` assembler and `yis` simulator can be downloaded from the links provided with Practical 2.

Further Reading

There are many descriptions of Y86-64 available on the web, and a description is included in Chapter 4 of the book below:

Computer Systems: A Programmer's Perspective
Randal E. Bryant and David R. O'Hallaron, Pearson, 2010.

A browser-based Y86-64 simulator which I modified from an existing Y86-32 one can be found here:

<http://www.cs.ox.ac.uk/people/alex.rogers/Y86-64/>

The unmodified version is available from here:

<http://xsnix.github.io/js-y86/>